

Lab2

June 20, 2023

0.1

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
```

```
[ ]: data = pd.read_csv("cars2.csv")
```

```
[ ]: data.head()
```

```
[ ]:      mileage      make  model    fuel   gear offerType  price    hp  year
0    235000      BMW    316   Diesel  Manual      Used   6800  116.0  2011
1     92800 Volkswagen   Golf  Gasoline  Manual      Used   6877  122.0  2011
2    149300      SEAT   Exeo  Gasoline  Manual      Used   6900  160.0  2011
3     96200   Renault  Megane  Gasoline  Manual      Used   6950  110.0  2011
4    156000   Peugeot    308  Gasoline  Manual      Used   6950  156.0  2011
```

```
[ ]: data_features = list(zip(
    [i for i in data.columns], #
    zip(
        [str(i) for i in data.dtypes], #
        [i for i in data.isnull().sum()] #
    )))
data_features #
```

```
[ ]: [('mileage', ('int64', 0)),
      ('make', ('object', 0)),
      ('model', ('object', 143)),
      ('fuel', ('object', 0)),
      ('gear', ('object', 182)),
      ('offerType', ('object', 0)),
      ('price', ('int64', 0)),
      ('hp', ('float64', 29)),
      ('year', ('int64', 0))]
```

0.2

```
[ ]: # ( )  
[(c, data[c].isnull().mean()) for c in data.columns]
```

```
[ ]: [('mileage', 0.0),  
      ('make', 0.0),  
      ('model', 0.0030815644865854973),  
      ('fuel', 0.0),  
      ('gear', 0.003921991164745178),  
      ('offerType', 0.0),  
      ('price', 0.0),  
      ('hp', 0.0006249326581187372),  
      ('year', 0.0)]
```

```
[ ]: #  
data.dropna(axis=1, how='any')
```

```
[ ]:      mileage      make      fuel      offerType  price  year  
0      235000      BMW      Diesel      Used      6800  2011  
1      92800  Volkswagen      Gasoline      Used      6877  2011  
2      149300      SEAT      Gasoline      Used      6900  2011  
3      96200      Renault      Gasoline      Used      6950  2011  
4      156000      Peugeot      Gasoline      Used      6950  2011  
...      ...      ...      ...      ...      ...  
46400      99      Fiat  Electric/Gasoline  Pre-registered  12990  2021  
46401      99      Fiat  Electric/Gasoline  Pre-registered  12990  2021  
46402      99      Fiat  Electric/Gasoline  Pre-registered  12990  2021  
46403      99      Fiat  Electric/Gasoline  Pre-registered  12990  2021  
46404      99      Fiat  Electric/Gasoline  Pre-registered  12990  2021
```

[46405 rows x 6 columns]

```
[ ]: # ( 50%)  
data.dropna(axis=1, thresh=23202)
```

```
[ ]:      mileage      make  model      fuel  gear      offerType  \  
0      235000      BMW      316      Diesel  Manual      Used  
1      92800  Volkswagen      Golf      Gasoline  Manual      Used  
2      149300      SEAT      Exeo      Gasoline  Manual      Used  
3      96200      Renault  Megane      Gasoline  Manual      Used  
4      156000      Peugeot      308      Gasoline  Manual      Used  
...      ...      ...      ...      ...      ...  
46400      99      Fiat      500  Electric/Gasoline  Manual  Pre-registered  
46401      99      Fiat      500  Electric/Gasoline  Manual  Pre-registered  
46402      99      Fiat      500  Electric/Gasoline  Manual  Pre-registered  
46403      99      Fiat      500  Electric/Gasoline  Manual  Pre-registered  
46404      99      Fiat      500  Electric/Gasoline  Manual  Pre-registered
```

	price	hp	year
0	6800	116.0	2011
1	6877	122.0	2011
2	6900	160.0	2011
3	6950	110.0	2011
4	6950	156.0	2011
...
46400	12990	71.0	2021
46401	12990	71.0	2021
46402	12990	71.0	2021
46403	12990	71.0	2021
46404	12990	71.0	2021

[46405 rows x 9 columns]

```
[ ]: # hp
def impute_na(df, variable, value):
    df[variable].fillna(value, inplace=True)
impute_na(data, 'hp', data['hp'].mean())
```

```
[ ]: # hp
data.isnull().sum()
```

```
[ ]: mileage      0
make            0
model          143
fuel           0
gear           182
offerType       0
price           0
hp             0
year           0
dtype: int64
```

0.3

```
[ ]: from sklearn.preprocessing import LabelEncoder
```

```
[ ]: le = LabelEncoder()
cat_enc_le = le.fit_transform(data['gear'])
```

```
[ ]: data['gear'].unique()
```

```
[ ]: array(['Manual', 'Automatic', nan, 'Semi-automatic'], dtype=object)
```

```
[ ]: np.unique(cat_enc_le)
```

```
[ ]: array([0, 1, 2, 3])
```

```
[ ]: le.inverse_transform([0, 1, 2, 3])
```

```
[ ]: array(['Automatic', 'Manual', 'Semi-automatic', nan], dtype=object)
```

```
[ ]: data['make'].unique()
```

```
[ ]: array(['BMW', 'Volkswagen', 'SEAT', 'Renault', 'Peugeot', 'Toyota',  
          'Opel', 'Mazda', 'Ford', 'Mercedes-Benz', 'Chevrolet', 'Audi',  
          'Fiat', 'Kia', 'Dacia', 'MINI', 'Hyundai', 'Skoda', 'Citroen',  
          'Infiniti', 'Suzuki', 'SsangYong', 'smart', 'Cupra', 'Volvo',  
          'Jaguar', 'Porsche', 'Nissan', 'Honda', 'Lada', 'Mitsubishi',  
          'Others', 'Lexus', 'Jeep', 'Maserati', 'Bentley', 'Land', 'Alfa',  
          'Subaru', 'Dodge', 'Microcar', 'Lamborghini', 'Baic', 'Tesla',  
          'Chrysler', '9ff', 'McLaren', 'Aston', 'Rolls-Royce', 'Alpine',  
          'Lancia', 'Abarth', 'DS', 'Daihatsu', 'Ligier', 'Ferrari',  
          'Caravans-Wohnm', 'Aixam', 'Piaggio', 'Zhidou', 'Morgan',  
          'Maybach', 'Tazzari', 'Trucks-Lkw', 'RAM', 'Iveco', 'DAF',  
          'Alpina', 'Polestar', 'Brilliance', 'FISKER', 'Cadillac',  
          'Trailer-Anhänger', 'Isuzu', 'Corvette', 'DFSK', 'Estrima'],  
          dtype=object)
```

```
[ ]: pip install category_encoders
```

```
Requirement already satisfied: category_encoders in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (2.6.1)  
Requirement already satisfied: numpy>=1.14.0 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (1.25.0)  
Requirement already satisfied: scikit-learn>=0.20.0 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (1.2.2)  
Requirement already satisfied: scipy>=1.0.0 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (1.10.1)  
Requirement already satisfied: statsmodels>=0.9.0 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (0.14.0)  
Requirement already satisfied: pandas>=1.0.5 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (2.0.2)  
Requirement already satisfied: patsy>=0.5.1 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
category_encoders) (0.5.3)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from  
pandas>=1.0.5->category_encoders) (2.8.2)
```

Requirement already satisfied: pytz>=2020.1 in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from
 pandas>=1.0.5->category_encoders) (2023.3)

Requirement already satisfied: tzdata>=2022.1 in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from
 pandas>=1.0.5->category_encoders) (2023.3)

Requirement already satisfied: six in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from
 patsy>=0.5.1->category_encoders) (1.16.0)

Requirement already satisfied: joblib>=1.1.1 in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from scikit-
 learn>=0.20.0->category_encoders) (1.2.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from scikit-
 learn>=0.20.0->category_encoders) (3.1.0)

Requirement already satisfied: packaging>=21.3 in
 /Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from
 statsmodels>=0.9.0->category_encoders) (23.1)

Note: you may need to restart the kernel to use updated packages.

```
[ ]: #CountEncoder
      from category_encoders.count import CountEncoder as ce_CountEncoder
```

```
[ ]: ce_CountEncoder1 = ce_CountEncoder()
      data_COUNT_ENC = ce_CountEncoder1.fit_transform(data[data.columns.
      ↪difference(['model']]))
```

```
[ ]: data_COUNT_ENC.head()
```

```
[ ]:      fuel    gear    hp  make  mileage  offerType  price  year
      0  15244  30380  116.0  2405   235000      40122   6800  2011
      1  28864  30380  122.0  6931    92800      40122   6877  2011
      2  28864  30380  160.0  1924   149300      40122   6900  2011
      3  28864  30380  110.0  2830    96200      40122   6950  2011
      4  28864  30380  156.0  1232   156000      40122   6950  2011
```

```
[ ]: data['offerType'].unique()
```

```
[ ]: array(['Used', 'Demonstration', "Employee's car", 'Pre-registered', 'New'],
      dtype=object)
```

```
[ ]: data_COUNT_ENC['offerType'].unique()
```

```
[ ]: array([40122, 2368, 1122, 2780, 13])
```

```
[ ]: ce_CountEncoder2 = ce_CountEncoder(normalize=True)
      data_FREQ_ENC = ce_CountEncoder2.fit_transform(data[data.columns.
      ↪difference(['model']]))
```

```
[ ]: data_FREQ_ENC['offerType'].unique()
```

```
[ ]: array([8.64605107e-01, 5.10289839e-02, 2.41784290e-02, 5.99073376e-02,
          2.80142226e-04])
```

```
[ ]: from category_encoders.helmert import HelmertEncoder as ce_HelmertEncoder
```

```
[ ]: #HelmertEncoder
ce_HelmertEncoder1 = ce_HelmertEncoder()
data_HELM_ENC = ce_HelmertEncoder1.fit_transform(data[data.columns.
↳ difference(['model'])], data['model'])
```

```
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-
packages/category_encoders/base_contrast_encoder.py:126: FutureWarning:
Intercept column might not be added anymore in future releases (c.f. issue #370)
  warnings.warn("Intercept column might not be added anymore in future releases
(c.f. issue #370)",
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-
packages/category_encoders/base_contrast_encoder.py:126: FutureWarning:
Intercept column might not be added anymore in future releases (c.f. issue #370)
  warnings.warn("Intercept column might not be added anymore in future releases
(c.f. issue #370)",
```

```
[ ]: data_HELM_ENC.head()
```

```
[ ]:   intercept  fuel_0  fuel_1  fuel_2  fuel_3  fuel_4  fuel_5  fuel_6  fuel_7  \
0           1    -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0
1           1     1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0
2           1     1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0
3           1     1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0
4           1     1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0   -1.0

   fuel_8  ...  make_73  make_74  make_75  mileage  offerType_0  offerType_1  \
0    -1.0  ...    -1.0    -1.0    -1.0   235000         -1.0         -1.0
1    -1.0  ...    -1.0    -1.0    -1.0    92800         -1.0         -1.0
2    -1.0  ...    -1.0    -1.0    -1.0   149300         -1.0         -1.0
3    -1.0  ...    -1.0    -1.0    -1.0    96200         -1.0         -1.0
4    -1.0  ...    -1.0    -1.0    -1.0   156000         -1.0         -1.0





   offerType_2  offerType_3  price  year
0         -1.0         -1.0   6800  2011
1         -1.0         -1.0   6877  2011
2         -1.0         -1.0   6900  2011
3         -1.0         -1.0   6950  2011
4         -1.0         -1.0   6950  2011
```

```
[5 rows x 98 columns]
```


0.4

```
[ ]: def diagnostic_plots(df, variable):  
    plt.figure(figsize=(15,6))  
    #  
    plt.subplot(1, 2, 1)  
    df[variable].hist(bins=30)  
    ## Q-Q plot  
    plt.subplot(1, 2, 2)  
    stats.probplot(df[variable], dist="norm", plot=plt)  
    plt.show()
```


```
[ ]: data.hist(figsize=(20,20))  
plt.show()
```





```
[ ]: diagnostic_plots(data, 'price')
```





```
[ ]: #
data['price'] = np.log(data['price'])
diagnostic_plots(data, 'price')
```





```
[ ]: #
data['price_reciprocal'] = 1 / (data['price'])
diagnostic_plots(data, 'price_reciprocal')
```



```
[ ]: #
data['price_sqr'] = data['price']**(1/2)
diagnostic_plots(data, 'price_sqr')
```





```
[ ]: #
data['price_exp1'] = data['price']**(1/1.5)
diagnostic_plots(data, 'price_exp1')
```



```
[ ]: data['price_exp2'] = data['price']**(2)
      diagnostic_plots(data, 'price_exp2')
```



```
[ ]: data['price_exp3'] = data['price']**(0.333)
      diagnostic_plots(data, 'price_exp3')
```



```
[ ]: #
data['price_boxcox'], param = stats.boxcox(data['price'])
print('          = {}'.format(param))
diagnostic_plots(data, 'price_boxcox')
```

= -1.8519809069200015

