

# Lab3

June 22, 2023

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
from sklearn.impute import SimpleImputer
from sklearn.impute import MissingIndicator
import scipy.stats as stats
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import RobustScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
```

```
[ ]: data = pd.read_csv("datasets/Accident.csv")
```

```
/var/folders/fs/5xh23h99763f_blp7m50x23h0000gq/T/ipykernel_68333/3503642161.py:1
: DtypeWarning: Columns (0) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
data = pd.read_csv("datasets/Accident.csv")
```

```
[ ]: data.head()
```

```
[ ]: Accident_Index 1st_Road_Class 1st_Road_Number 2nd_Road_Class \
0 200501BS00001 A 3218 NaN
1 200501BS00002 B 450 C
2 200501BS00003 C 0 NaN
3 200501BS00004 A 3220 NaN
4 200501BS00005 Unclassified 0 NaN

2nd_Road_Number Accident_Severity Carriageway_Hazards Date \
0 0.0 Serious NaN 2005-01-04
1 0.0 Slight NaN 2005-01-05
2 0.0 Slight NaN 2005-01-06
3 0.0 Slight NaN 2005-01-07
4 0.0 Slight NaN 2005-01-10
```

	Day_of_Week	Did_Police_Officer_Attend_Scene_of_Accident	...	\
0	Tuesday	1.0	...	
1	Wednesday	1.0	...	
2	Thursday	1.0	...	
3	Friday	1.0	...	
4	Monday	1.0	...	

	Police_Force	Road_Surface_Conditions	Road_Type	\
0	Metropolitan Police	Wet or damp	Single carriageway	
1	Metropolitan Police	Dry	Dual carriageway	
2	Metropolitan Police	Dry	Single carriageway	
3	Metropolitan Police	Dry	Single carriageway	
4	Metropolitan Police	Wet or damp	Single carriageway	

	Special_Conditions_at_Site	Speed_limit	Time	Urban_or_Rural_Area	\
0	NaN	30	17:42	Urban	
1	NaN	30	17:36	Urban	
2	NaN	30	00:15	Urban	
3	NaN	30	10:35	Urban	
4	NaN	30	21:13	Urban	

	Weather_Conditions	Year	InScotland
0	Raining no high winds	2005	No
1	Fine no high winds	2005	No
2	Fine no high winds	2005	No
3	Fine no high winds	2005	No
4	Fine no high winds	2005	No

[5 rows x 34 columns]

```
[ ]: # ( 25%)
data.dropna(axis=1, thresh=37625)
```

	Accident_Index	1st_Road_Class	1st_Road_Number	2nd_Road_Number	\
0	200501BS00001	A	3218	0.0	
1	200501BS00002	B	450	0.0	
2	200501BS00003	C	0	0.0	
3	200501BS00004	A	3220	0.0	
4	200501BS00005	Unclassified	0	0.0	
...	...	...	...	...	
50164	2005070502900	Unclassified	0	0.0	
50165	2005070502901	Unclassified	0	0.0	
50166	2005070502902	A	537	0.0	
50167	2005070502903	B	5210	574.0	
50168	2005070502904	A	527	0.0	

	Accident_Severity	Date	Day_of_Week	\
0	Serious	2005-01-04	Tuesday	
1	Slight	2005-01-05	Wednesday	
2	Slight	2005-01-06	Thursday	
3	Slight	2005-01-07	Friday	
4	Slight	2005-01-10	Monday	
...	...	...	...	
50164	Serious	2005-07-03	Sunday	
50165	Slight	2005-07-03	Sunday	
50166	Serious	2005-07-03	Sunday	
50167	Slight	2005-07-04	Monday	
50168	Slight	2005-07-04	Monday	

	Did_Police_Officer_Attend_Scene_of_Accident	\
0	1.0	
1	1.0	
2	1.0	
3	1.0	
4	1.0	
...	...	
50164	1.0	
50165	1.0	
50166	1.0	
50167	2.0	
50168	1.0	

	Junction_Control	Junction_Detail	...	\
0	Data missing or out of range	Not at junction or within 20 metres	...	
1	Auto traffic signal	Crossroads	...	
2	Data missing or out of range	Not at junction or within 20 metres	...	
3	Data missing or out of range	Not at junction or within 20 metres	...	
4	Data missing or out of range	Not at junction or within 20 metres	...	
...	...	...	...	
50164	Give way or uncontrolled	Crossroads	...	
50165	Give way or uncontrolled	Crossroads	...	
50166	Data missing or out of range	Not at junction or within 20 metres	...	
50167	Give way or uncontrolled	Roundabout	...	
50168	Give way or uncontrolled	Private drive or entrance	...	

	Pedestrian_Crossing-Physical_Facilities	Police_Force	\
0	1.0	Metropolitan Police	
1	5.0	Metropolitan Police	
2	0.0	Metropolitan Police	
3	0.0	Metropolitan Police	
4	0.0	Metropolitan Police	
...	...	...	
50164	0.0	Cheshire	

50165	0.0	Cheshire
50166	0.0	Cheshire
50167	0.0	Cheshire
50168	0.0	Cheshire

	Road_Surface_Conditions	Road_Type	Speed_limit	Time \
0	Wet or damp	Single carriageway	30	17:42
1	Dry	Dual carriageway	30	17:36
2	Dry	Single carriageway	30	00:15
3	Dry	Single carriageway	30	10:35
4	Wet or damp	Single carriageway	30	21:13
...	...	...	...	...
50164	Dry	Single carriageway	40	18:29
50165	Dry	Single carriageway	30	20:07
50166	Dry	Single carriageway	60	16:04
50167	Wet or damp	Roundabout	30	16:45
50168	Wet or damp	Single carriageway	30	13:35

	Urban_or_Rural_Area	Weather_Conditions	Year	InScotland
0	Urban	Raining no high winds	2005	No
1	Urban	Fine no high winds	2005	No
2	Urban	Fine no high winds	2005	No
3	Urban	Fine no high winds	2005	No
4	Urban	Fine no high winds	2005	No
...	...	...	...	...
50164	Urban	Fine no high winds	2005	No
50165	Urban	Fine no high winds	2005	No
50166	Rural	Fine no high winds	2005	No
50167	Urban	Fine no high winds	2005	No
50168	Urban	Fine no high winds	2005	No

[50169 rows x 31 columns]

```
[ ]: #
def impute_na(df, variable, value):
    df[variable].fillna(value, inplace=True)
impute_na(data, 'Number_of_Vehicles', data['Number_of_Vehicles'].mean())
```

```
[ ]: data.describe()
```

	1st_Road_Number	2nd_Road_Number \
count	50169.000000	50072.000000
mean	865.940920	298.751578
std	1704.128271	1100.096090
min	0.000000	0.000000
25%	0.000000	0.000000
50%	61.000000	0.000000

75%	574.000000	0.000000
max	7076.000000	8228.000000

	Did_Police_Officer_Attend_Scene_of_Accident	Latitude \
count	50159.000000	50156.000000
mean	1.205188	52.491185
std	0.429536	1.079310
min	1.000000	51.289060
25%	1.000000	51.509316
50%	1.000000	51.608980
75%	1.000000	53.498600
max	3.000000	55.047995

	Location_Easting_OSGR	Location_Northing_OSGR	Longitude \
count	50156.000000	50156.000000	50156.000000
mean	451681.899673	289155.002193	-1.269933
std	86087.912299	119358.039886	1.256771
min	296300.000000	156110.000000	-3.602717
25%	367370.000000	180480.000000	-2.494029
50%	513030.000000	191740.000000	-0.372561
75%	531292.500000	400320.000000	-0.110122
max	559570.000000	573030.000000	0.300802

	Number_of_Casualties	Number_of_Vehicles \
count	50169.000000	50169.000000
mean	1.302338	1.827403
std	0.743737	0.686198
min	1.000000	1.000000
25%	1.000000	1.000000
50%	1.000000	2.000000
75%	1.000000	2.000000
max	23.000000	18.000000

	Pedestrian_Crossing-Human_Control \
count	50156.000000
mean	0.006579
std	0.103746
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	2.000000

	Pedestrian_Crossing-Physical_Facilities	Speed_limit	Year
count	50159.000000	50169.000000	50169.0
mean	0.905162	33.573521	2005.0
std	1.914236	9.766182	0.0

min	0.000000	10.000000	2005.0
25%	0.000000	30.000000	2005.0
50%	0.000000	30.000000	2005.0
75%	0.000000	30.000000	2005.0
max	8.000000	70.000000	2005.0

```
[ ]: def obj_col(column):
      return column[1] == 'object'

col_names = []
for col in list(filter(obj_col, list(zip(list(data.columns), list(data.
    dtypes))))):
    col_names.append(col[0])
col_names.append('Speed_limit')
```

```
[ ]: X_ALL = data.drop(col_names, axis=1)
```

```
[ ]: #
      #
      def arr_to_df(arr_scaled):
          res = pd.DataFrame(arr_scaled, columns=X_ALL.columns)
          return res
```

## 0.1 StandardScaler

```
[ ]: #
      X_train, X_test, y_train, y_test = train_test_split(X_ALL, data['Speed_limit'],
                                                          test_size=0.2,
                                                          random_state=1)

      # DataFrame
      X_train_df = arr_to_df(X_train)
      X_test_df = arr_to_df(X_test)

      X_train_df.shape, X_test_df.shape
```

```
[ ]: ((40135, 12), (10034, 12))
```

```
[ ]: # StandardScaler
      cs11 = StandardScaler()
      data_cs11_scaled_temp = cs11.fit_transform(X_ALL)
      # DataFrame
      data_cs11_scaled = arr_to_df(data_cs11_scaled_temp)
      data_cs11_scaled
```

```
[ ]:      1st_Road_Number  2nd_Road_Number  \
0          1.380226      -0.271571
1         -0.244081      -0.271571
```

2	-0.508148	-0.271571
3	1.381400	-0.271571
4	-0.508148	-0.271571
...	...	...
50164	-0.508148	-0.271571
50165	-0.508148	-0.271571
50166	-0.193028	-0.271571
50167	2.549164	0.250206
50168	-0.198896	-0.271571

	Did_Police_Officer_Attend_Scene_of_Accident	Latitude	\
0		-0.477700	-0.928463
1		-0.477700	-0.899760
2		-0.477700	-0.894918
3		-0.477700	-0.934628
4		-0.477700	-0.922296
...		...	...
50164		-0.477700	0.638403
50165		-0.477700	0.556717
50166		-0.477700	0.709095
50167		1.850416	0.854118
50168		-0.477700	0.609283

	Location_Easting_OSGR	Location_Northing_OSGR	Longitude	\
0	0.859573	-0.929272	0.858369	
1	0.842033	-0.900702	0.842027	
2	0.846098	-0.895759	0.846205	
3	0.873745	-0.935221	0.872141	
4	0.887219	-0.922570	0.885862	
...	...	...	...	
50164	-0.936400	0.634442	-0.926521	
50165	-0.951733	0.552335	-0.941498	
50166	-0.616375	0.704819	-0.598538	
50167	-1.018177	0.851857	-1.012682	
50168	-0.736021	0.604532	-0.721037	

	Number_of_Casualties	Number_of_Vehicles	\
0	-0.406516	-1.205791	
1	-0.406516	-1.205791	
2	-0.406516	0.251529	
3	-0.406516	-1.205791	
4	-0.406516	-1.205791	
...	...	...	
50164	6.316361	1.708849	
50165	0.938059	0.251529	
50166	-0.406516	-1.205791	
50167	0.938059	0.251529	

```
50168          -0.406516          0.251529
```

```

Pedestrian_Crossing-Human_Control  \
0          -0.06342
1          -0.06342
2          -0.06342
3          -0.06342
4          -0.06342
...
50164          -0.06342
50165          -0.06342
50166          -0.06342
50167          -0.06342
50168          -0.06342

```

```

Pedestrian_Crossing-Physical_Facilities  Year
0          0.049544  0.0
1          2.139172  0.0
2         -0.472863  0.0
3         -0.472863  0.0
4         -0.472863  0.0
...
50164         -0.472863  0.0
50165         -0.472863  0.0
50166         -0.472863  0.0
50167         -0.472863  0.0
50168         -0.472863  0.0

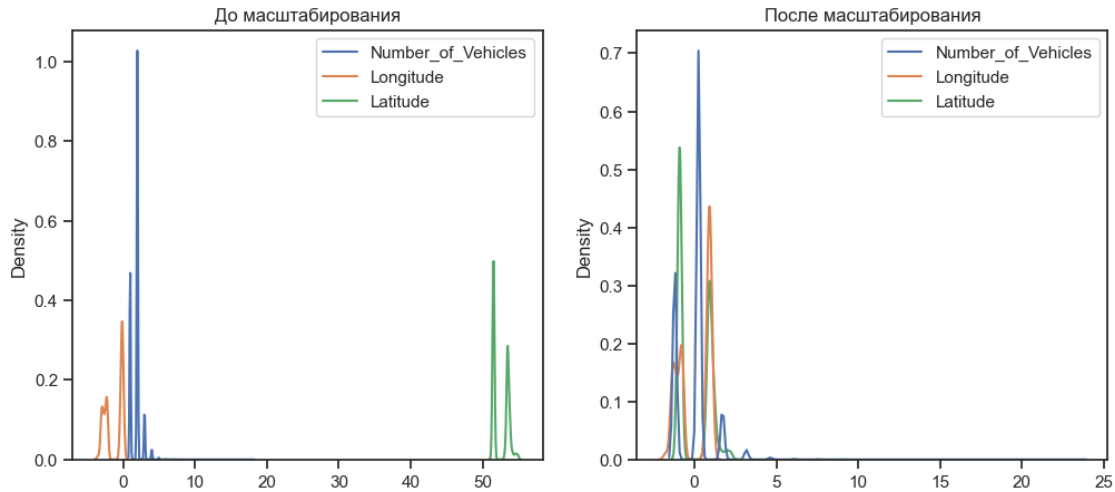
```

```
[50169 rows x 12 columns]
```

```
[ ]: #
def draw_kde(col_list, df1, df2, label1, label2):
    fig, (ax1, ax2) = plt.subplots(
        ncols=2, figsize=(12, 5))
    #
    ax1.set_title(label1)
    sns.kdeplot(data=df1[col_list], ax=ax1)
    #
    ax2.set_title(label2)
    sns.kdeplot(data=df2[col_list], ax=ax2)
    plt.show()
```

```
[ ]: draw_kde(['Number_of_Vehicles', 'Longitude', 'Latitude'], data,
    ↪data_cs11_scaled, ' ', ' ')
```





## 0.2 “Mean Normalisation”

```
[ ]: #
X_train, X_test, y_train, y_test = train_test_split(X_ALL, data['Speed_limit'],
                                                    test_size=0.2,
                                                    random_state=1)

#           DataFrame
X_train_df = arr_to_df(X_train)
X_test_df = arr_to_df(X_test)

X_train_df.shape, X_test_df.shape
```

```
[ ]: ((40135, 12), (10034, 12))
```

```
[ ]: class MeanNormalisation:

    def fit(self, param_df):
        self.means = X_train.mean(axis=0)
        maxs = X_train.max(axis=0)
        mins = X_train.min(axis=0)
        self.ranges = maxs - mins

    def transform(self, param_df):
        param_df_scaled = (param_df - self.means) / self.ranges
        return param_df_scaled

    def fit_transform(self, param_df):
        self.fit(param_df)
        return self.transform(param_df)
```

```
[ ]: sc21 = MeanNormalisation()
data_cs21_scaled = sc21.fit_transform(X_ALL)
data_cs21_scaled.describe()
```

```
[ ]:      1st_Road_Number  2nd_Road_Number  \
count      50169.000000      50072.000000
mean        -0.000535        -0.000135
std          0.240832         0.133702
min         -0.122912        -0.036444
25%         -0.122912        -0.036444
50%         -0.114292        -0.036444
75%         -0.041793        -0.036444
max          0.877088         0.963556
```

```
      Did_Police_Officer_Attend_Scene_of_Accident  Latitude  \
count      50159.000000      50156.000000
mean        -0.000419        -0.001192
std          0.214768         0.287132
min         -0.103013        -0.320997
25%         -0.103013        -0.262402
50%         -0.103013        -0.235888
75%         -0.103013         0.266813
max          0.896987         0.679003
```

```
      Location_Easting_OSGR  Location_Northing_OSGR  Longitude  \
count      50156.000000      50156.000000      50156.000000
mean          0.001293        -0.001188         0.001270
std          0.327592         0.286285         0.322451
min         -0.589985        -0.320302        -0.597255
25%         -0.319541        -0.261849        -0.312798
50%          0.234742        -0.234842         0.231509
75%          0.304236         0.265446         0.298844
max          0.411841         0.679698         0.404275
```

```
      Number_of_Casualties  Number_of_Vehicles  \
count      50169.000000      50169.000000
mean        -0.000113         0.000170
std          0.033806         0.040365
min         -0.013856        -0.048501
25%         -0.013856        -0.048501
50%         -0.013856         0.010323
75%         -0.013856         0.010323
max          0.986144         0.951499
```

```
      Pedestrian_Crossing-Human_Control  \
count      50156.000000
mean        -0.000025
```

std	0.051873
min	-0.003315
25%	-0.003315
50%	-0.003315
75%	-0.003315
max	0.996685

	Pedestrian_Crossing-Physical_Facilities	Year
count	50159.000000	0.0
mean	0.000120	NaN
std	0.239279	NaN
min	-0.113026	NaN
25%	-0.113026	NaN
50%	-0.113026	NaN
75%	-0.113026	NaN
max	0.886974	NaN

```
[ ]: cs22 = MeanNormalisation()
cs22.fit(X_train)
data_cs22_scaled_train = cs22.transform(X_train)
data_cs22_scaled_test = cs22.transform(X_test)
```

```
[ ]: data_cs22_scaled_train.describe()
```

```
[ ]:      1st_Road_Number  2nd_Road_Number  \
count      4.013500e+04      4.005800e+04
mean      4.558733e-18      -1.529889e-18
std       2.416528e-01      1.340982e-01
min      -1.229122e-01      -3.644368e-02
25%      -1.229122e-01      -3.644368e-02
50%      -1.142916e-01      -3.644368e-02
75%      -4.151033e-02      -3.644368e-02
max       8.770878e-01      9.635563e-01
```

	Did_Police_Officer_Attend_Scene_of_Accident	Latitude	\
count	4.012600e+04	4.012400e+04	
mean	-3.652231e-17	2.606982e-15	
std	2.154150e-01	2.873328e-01	
min	-1.030130e-01	-3.209971e-01	
25%	-1.030130e-01	-2.623026e-01	
50%	-1.030130e-01	-2.354915e-01	
75%	-1.030130e-01	2.672322e-01	
max	8.969870e-01	6.790029e-01	

	Location_Easting_OSGR	Location_Northing_OSGR	Longitude	\
count	4.012400e+04	4.012400e+04	4.012400e+04	
mean	6.330850e-18	-4.830040e-17	5.954541e-17	

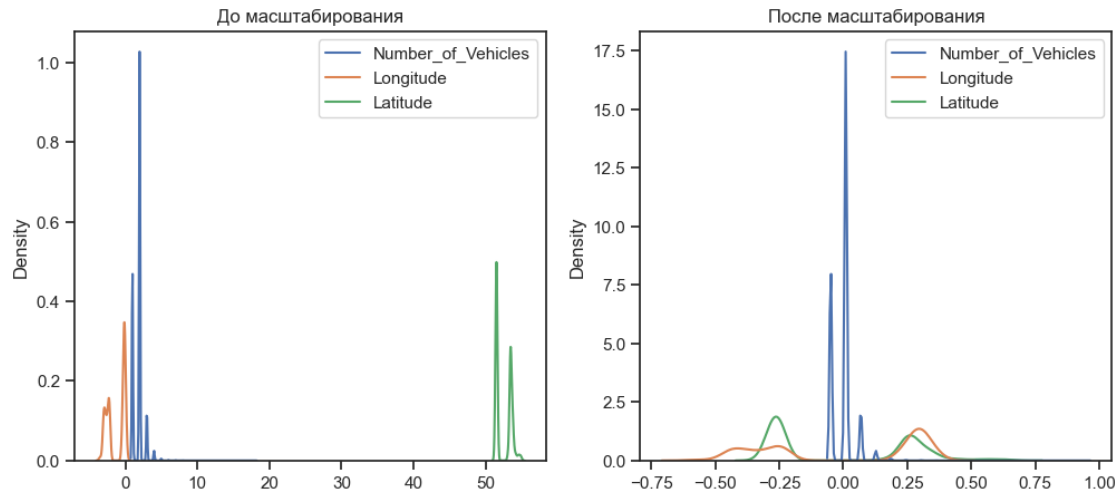
std	3.275957e-01	2.864881e-01	3.224602e-01
min	-5.881587e-01	-3.203017e-01	-5.957253e-01
25%	-3.206257e-01	-2.616573e-01	-3.139790e-01
50%	2.340187e-01	-2.345298e-01	2.308289e-01
75%	3.039605e-01	2.658540e-01	2.985749e-01
max	4.118413e-01	6.796983e-01	4.042747e-01

	Number_of_Casualties	Number_of_Vehicles \
count	4.013500e+04	4.013500e+04
mean	3.208817e-18	5.399664e-18
std	3.400279e-02	4.039157e-02
min	-1.385551e-02	-4.850101e-02
25%	-1.385551e-02	-4.850101e-02
50%	-1.385551e-02	1.032251e-02
75%	-1.385551e-02	1.032251e-02
max	9.861445e-01	9.514990e-01

	Pedestrian_Crossing-Human_Control \
count	4.012300e+04
mean	-4.427278e-19
std	5.225562e-02
min	-3.314807e-03
25%	-3.314807e-03
50%	-3.314807e-03
75%	-3.314807e-03
max	9.966852e-01


	Pedestrian_Crossing-Physical_Facilities	Year
count	4.012700e+04	0.0
mean	-9.296357e-19	NaN
std	2.388316e-01	NaN
min	-1.130255e-01	NaN
25%	-1.130255e-01	NaN
50%	-1.130255e-01	NaN
75%	-1.130255e-01	NaN
max	8.869745e-01	NaN

```
[ ]: draw_kde(['Number_of_Vehicles', 'Longitude', 'Latitude'], data,
↳data_cs21_scaled, ' ', ' ')
```



```
[ ]: draw_kde(['Number_of_Vehicles', 'Longitude', 'Latitude'],
↳data_cs22_scaled_train, data_cs22_scaled_test, ' ', ' ')

```



### 0.3 MinMax-

```
[ ]: # StandardScaler
cs31 = MinMaxScaler()
data_cs31_scaled_temp = cs31.fit_transform(X_ALL)
# DataFrame
data_cs31_scaled = arr_to_df(data_cs31_scaled_temp)
data_cs31_scaled.describe()

```

```

[ ]:      1st_Road_Number  2nd_Road_Number  \
count      50169.000000      50072.000000
mean         0.122377         0.036309
std          0.240832         0.133702
min          0.000000         0.000000
25%          0.000000         0.000000
50%          0.008621         0.000000
75%          0.081119         0.000000
max          1.000000         1.000000

      Did_Police_Officer_Attend_Scene_of_Accident  Latitude  \
count                        50159.000000  50156.000000
mean                          0.102594    0.319805
std                            0.214768    0.287132
min                            0.000000    0.000000
25%                            0.000000    0.058595
50%                            0.000000    0.085109
75%                            0.000000    0.587810
max                            1.000000    1.000000

      Location_Easting_OSGR  Location_Northing_OSGR  Longitude  \
count      50156.000000      50156.000000  50156.000000
mean         0.590200         0.319114    0.597610
std          0.326995         0.286285    0.321959
min          0.000000         0.000000    0.000000
25%          0.269951         0.058452    0.284023
50%          0.823223         0.085460    0.827498
75%          0.892591         0.585748    0.894730
max          1.000000         1.000000    1.000000

      Number_of_Casualties  Number_of_Vehicles  \
count      50169.000000      50169.000000
mean         0.013743         0.048671
std          0.033806         0.040365
min          0.000000         0.000000
25%          0.000000         0.000000
50%          0.000000         0.058824
75%          0.000000         0.058824
max          1.000000         1.000000

      Pedestrian_Crossing-Human_Control  \
count                        50156.000000
mean                          0.003290
std                            0.051873
min                            0.000000
25%                            0.000000
50%                            0.000000

```

```

75%                0.000000
max                1.000000

```

	Pedestrian_Crossing-Physical_Facilities	Year
count	50159.000000	50169.0
mean	0.113145	0.0
std	0.239279	0.0
min	0.000000	0.0
25%	0.000000	0.0
50%	0.000000	0.0
75%	0.000000	0.0
max	1.000000	0.0

```

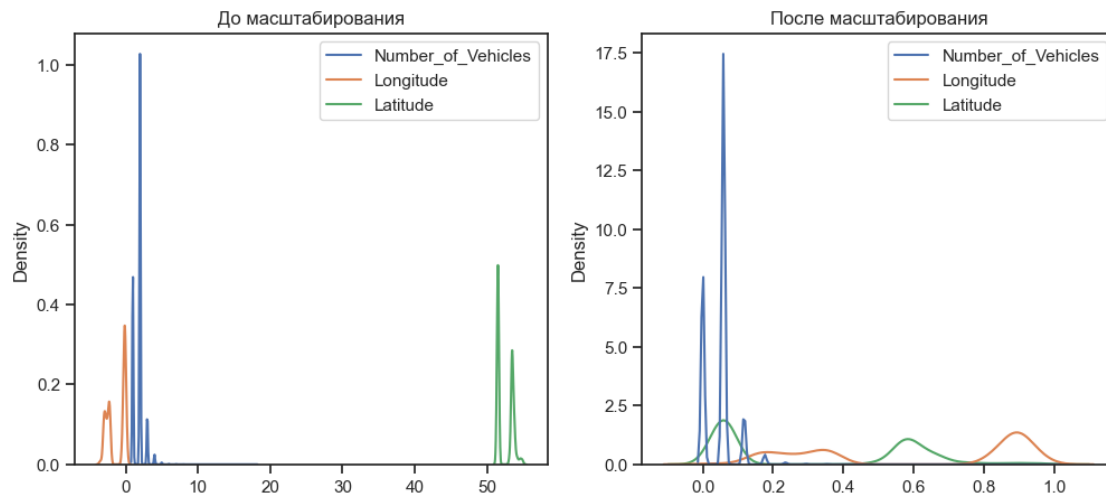
[ ]: cs32 = MinMaxScaler()
cs32.fit(X_train)
data_cs32_scaled_train_temp = cs32.transform(X_train)
data_cs32_scaled_test_temp = cs32.transform(X_test)
# DataFrame
data_cs32_scaled_train = arr_to_df(data_cs32_scaled_train_temp)
data_cs32_scaled_test = arr_to_df(data_cs32_scaled_test_temp)

```

```

[ ]: draw_kde(['Number_of_Vehicles', 'Longitude', 'Latitude'], data,
↳data_cs31_scaled, ' ', ' ')

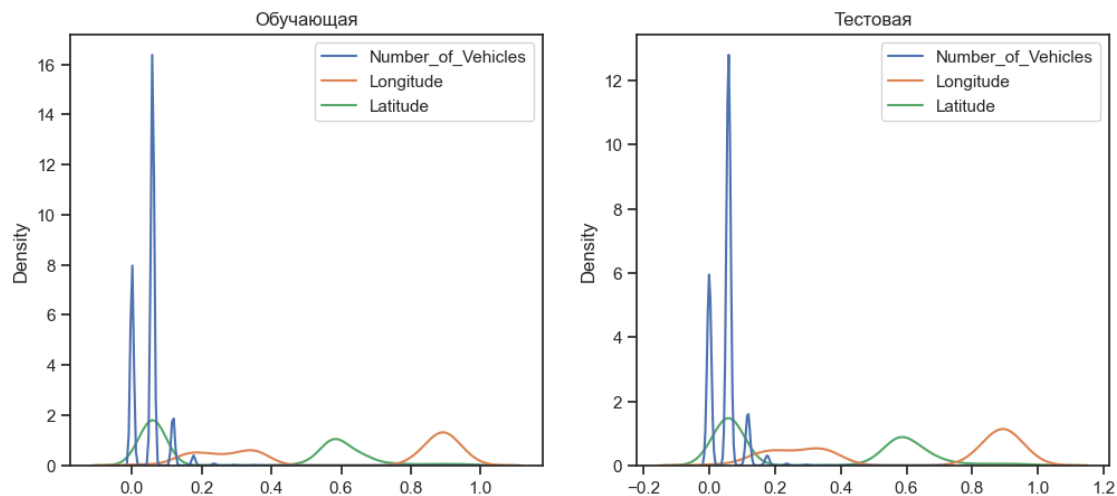
```



```

[ ]: draw_kde(['Number_of_Vehicles', 'Longitude', 'Latitude'],
↳data_cs32_scaled_train, data_cs32_scaled_test, ' ', ' ')

```



## 0.4

```
[ ]: data2 = pd.read_csv("datasets/Car_Sales.csv")
```

```
[ ]: data2.head()
```

```
[ ]:
Manufacturer  Model  Sales_in_thousands  __year_resale_value  Vehicle_type \
0      Acura  Integra          16.919             16.360      Passenger
1      Acura    TL          39.384             19.875      Passenger
2      Acura    CL          14.114             18.225      Passenger
3      Acura    RL           8.588             29.725      Passenger
4      Audi     A4          20.397             22.255      Passenger
```

```

Price_in_thousands  Engine_size  Horsepower  Wheelbase  Width  Length  \
0              21.50           1.8      140.0      101.2   67.3   172.4
1              28.40           3.2      225.0      108.1   70.3   192.9
2              NaN           3.2      225.0      106.9   70.6   192.0
3              42.00           3.5      210.0      114.6   71.4   196.6
4              23.99           1.8      150.0      102.6   68.2   178.0
```

```

Curb_weight  Fuel_capacity  Fuel_efficiency  Latest_Launch  \
0      2.639           13.2           28.0      2/2/2012
1      3.517           17.2           25.0      6/3/2011
2      3.470           17.2           26.0      1/4/2012
3      3.850           18.0           22.0      3/10/2011
4      2.998           16.4           27.0      10/8/2011
```

```

Power_perf_factor
0      58.280150
1      91.370778
```



```

2          NaN
3      91.389779
4      62.777639

```

```
[ ]: data2.describe()
```

```
[ ]:
      Sales_in_thousands  __year_resale_value  Price_in_thousands  \
count          157.000000          121.000000          155.000000
mean           52.998076           18.072975           27.390755
std            68.029422           11.453384           14.351653
min             0.110000            5.160000            9.235000
25%            14.114000           11.260000           18.017500
50%            29.450000           14.180000           22.799000
75%            67.956000           19.875000           31.947500
max           540.561000           67.550000           85.500000

```

```

      Engine_size  Horsepower  Wheelbase  Width  Length  \
count    156.000000  156.000000  156.000000  156.000000  156.000000
mean       3.060897  185.948718  107.487179   71.150000  187.343590
std        1.044653   56.700321    7.641303    3.451872   13.431754
min         1.000000   55.000000   92.600000   62.600000  149.400000
25%         2.300000  149.500000  103.000000   68.400000  177.575000
50%         3.000000  177.500000  107.000000   70.550000  187.900000
75%         3.575000  215.000000  112.200000   73.425000  196.125000
max         8.000000  450.000000  138.700000   79.900000  224.500000

```

```

      Curb_weight  Fuel_capacity  Fuel_efficiency  Power_perf_factor
count    155.000000    156.000000    154.000000    155.000000
mean       3.378026    17.951923     23.844156     77.043591
std        0.630502     3.887921     4.282706     25.142664
min        1.895000    10.300000    15.000000     23.276272
25%        2.971000    15.800000    21.000000     60.407707
50%        3.342000    17.200000    24.000000     72.030917
75%        3.799500    19.575000    26.000000     89.414878
max        5.572000    32.000000    45.000000    188.144323

```

```
[ ]: def diagnostic_plots(df, variable, title):
      fig, ax = plt.subplots(figsize=(10,7))
      #
      plt.subplot(2, 2, 1)
      df[variable].hist(bins=30)
      ## Q-Q plot
      plt.subplot(2, 2, 2)
      stats.probplot(df[variable], dist="norm", plot=plt)
      # violinplot
      plt.subplot(2, 2, 3)
      sns.violinplot(x=df[variable])


```

```
# boxplot
plt.subplot(2, 2, 4)
sns.boxplot(x=df[variable])
fig.suptitle(title)
plt.show()
```

```
[ ]: diagnostic_plots(data2, 'Price_in_thousands', 'Price_in_thousands - original')
```

/var/folders/fs/5xh23h99763f\_blp7m50x23h0000gq/T/ipykernel\_68333/4201870494.py:4  
: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(2, 2, 1)
```




```
[ ]: diagnostic_plots(data2, 'Engine_size', 'Engine_size - original')
```

/var/folders/fs/5xh23h99763f\_blp7m50x23h0000gq/T/ipykernel\_68333/4201870494.py:4  
: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call

```
ax.remove() as needed.
plt.subplot(2, 2, 1)
```

Engine\_size - original



```
[ ]: #
from enum import Enum
class OutlierBoundaryType(Enum):
    SIGMA = 1
    QUANTILE = 2
    IRQ = 3

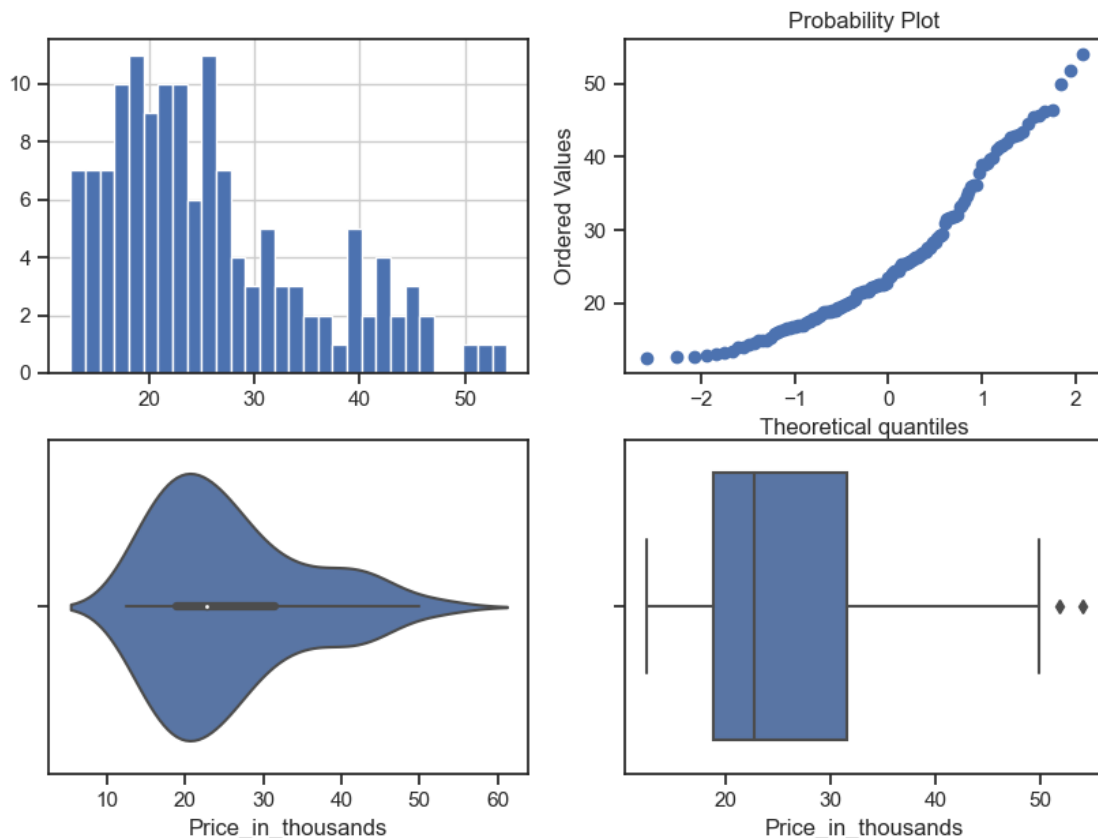
[ ]: #
def get_outlier_boundaries(df, col):
    lower_boundary = df[col].quantile(0.05)
    upper_boundary = df[col].quantile(0.95)
    return lower_boundary, upper_boundary
```

0.5 (number\_of\_reviews)

```
[ ]: #
lower_boundary, upper_boundary = get_outlier_boundaries(data2,
    ↪ "Price_in_thousands")
#
outliers_temp = np.where(data2["Price_in_thousands"] > upper_boundary, True,
    ↪ np.where(data2["Price_in_thousands"] < lower_boundary,
    ↪ True, False))
#
data_trimmed = data2.loc[~(outliers_temp), ]
title = ' -{ }, -{ }, -{ }'.format("Price_in_thousands", "QUANTILE",
    ↪ data_trimmed.shape[0])
diagnostic_plots(data_trimmed, "Price_in_thousands", title)
```

```
/var/folders/fs/5xh23h99763f_blp7m50x23h0000gq/T/ipykernel_68333/4201870494.py:4
: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated
since 3.6 and will be removed two minor releases later; explicitly call
ax.remove() as needed.
plt.subplot(2, 2, 1)
```

Поле-Price\_in\_thousands, метод-QUANTILE, строк-141



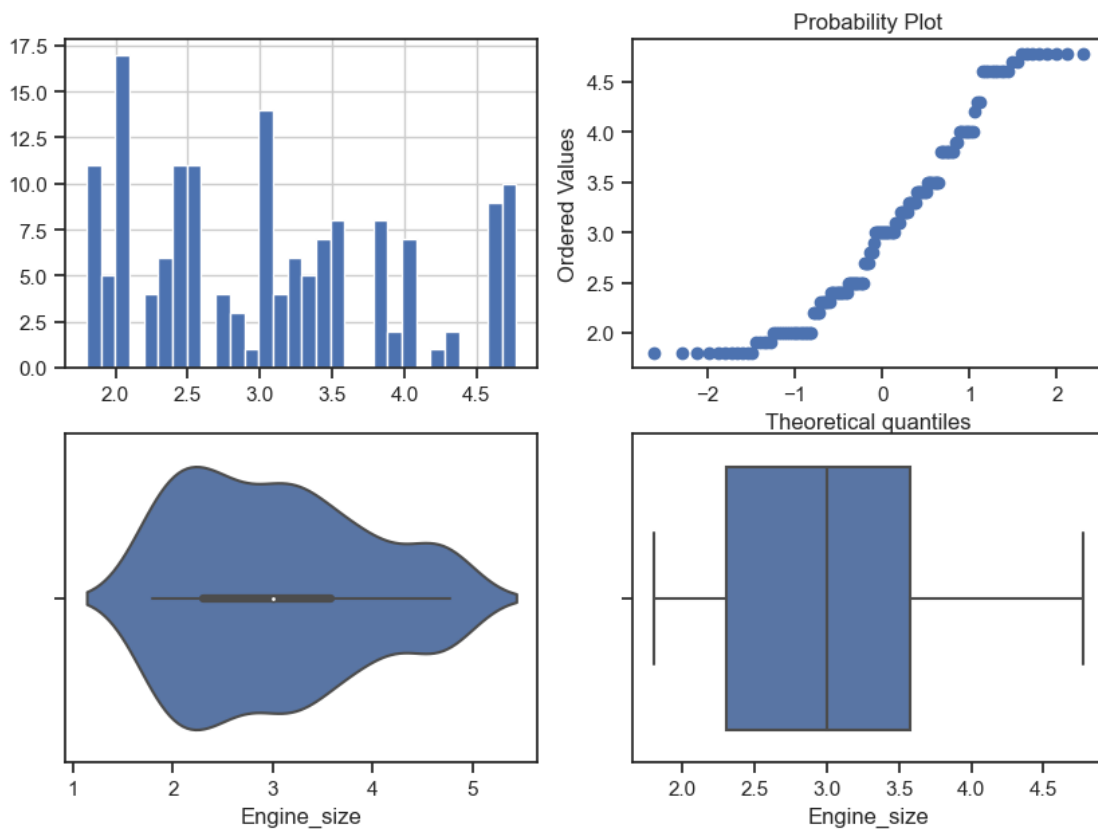
## 0.6

```
[ ]: #
lower_boundary, upper_boundary = get_outlier_boundaries(data2, "Engine_size")
#
data2["Engine_size"] = np.where(data2["Engine_size"] > upper_boundary,
    ↪upper_boundary,
    np.where(data2["Engine_size"] < lower_boundary,
    ↪lower_boundary, data2["Engine_size"]))
title = ' -{ }, -{ }'.format("Engine_size", "QUANTILE")
diagnostic_plots(data2, "Engine_size", title)
```

/var/folders/fs/5xh23h99763f\_blp7m50x23h0000gq/T/ipykernel\_68333/4201870494.py:4  
: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(2, 2, 1)
```

Поле-Engine\_size, метод-QUANTILE



## 0.7

```
[ ]: data2.dtypes
```

```
[ ]: Manufacturer      object
      Model            object
      Sales_in_thousands float64
      __year_resale_value float64
      Vehicle_type     object
      Price_in_thousands float64
      Engine_size      float64
      Horsepower       float64
      Wheelbase        float64
      Width            float64
      Length           float64
      Curb_weight      float64
      Fuel_capacity     float64
      Fuel_efficiency   float64
      Latest_Launch    object
      Power_perf_factor float64
      dtype: object
```

```
[ ]: #
      data2["Date"] = data2.apply(lambda x: pd.to_datetime(x["Latest_Launch"],
      ↪format='%m/%d/%Y'), axis=1)
```

```
[ ]: data2.head(5)
```

```
[ ]:  Manufacturer      Model  Sales_in_thousands  __year_resale_value  Vehicle_type  \
0      Acura      Integra           16.919           16.360      Passenger
1      Acura         TL           39.384           19.875      Passenger
2      Acura         CL           14.114           18.225      Passenger
3      Acura         RL            8.588           29.725      Passenger
4      Audi         A4           20.397           22.255      Passenger

      Price_in_thousands  Engine_size  Horsepower  Wheelbase  Width  Length  \
0              21.50           1.8         140.0        101.2    67.3    172.4
1              28.40           3.2         225.0        108.1    70.3    192.9
2               NaN           3.2         225.0        106.9    70.6    192.0
3              42.00           3.5         210.0        114.6    71.4    196.6
4              23.99           1.8         150.0        102.6    68.2    178.0

      Curb_weight  Fuel_capacity  Fuel_efficiency  Latest_Launch  \
0          2.639         13.2           28.0        2/2/2012
1          3.517         17.2           25.0        6/3/2011
2          3.470         17.2           26.0        1/4/2012
3          3.850         18.0           22.0        3/10/2011
4          2.998         16.4           27.0       10/8/2011
```

	Power_perf_factor	Date
0	58.280150	2012-02-02
1	91.370778	2011-06-03
2	NaN	2012-01-04
3	91.389779	2011-03-10
4	62.777639	2011-10-08

0.8

0.9 ( )

```
[ ]: data3 = pd.read_csv("datasets/Marketing.csv")
```

```
[ ]: data3.head()
```

```
[ ]:
   Income  Kidhome  Teenhome  Recency  MntWines  MntFruits  MntMeatProducts  \
0  58138.0        0         0       58       635         88           546
1  46344.0        1         1       38        11          1            6
2  71613.0        0         0       26       426         49          127
3  26646.0        1         0       26        11          4           20
4  58293.0        1         0       94       173         43          118
```

	MntFishProducts	MntSweetProducts	MntGoldProds	...	marital_Together	\
0	172		88	88 ...	0	
1	2		1	6 ...	0	
2	111		21	42 ...	1	
3	10		3	5 ...	1	
4	46		27	15 ...	0	

	marital_Widow	education_2n Cycle	education_Basic	education_Graduation	\
0	0	0	0		1
1	0	0	0		1
2	0	0	0		1
3	0	0	0		1
4	0	0	0		0

	education_Master	education_PhD	MntTotal	MntRegularProds	\
0	0	0	1529	1441	
1	0	0	21	15	
2	0	0	734	692	
3	0	0	48	43	
4	0	1	407	392	

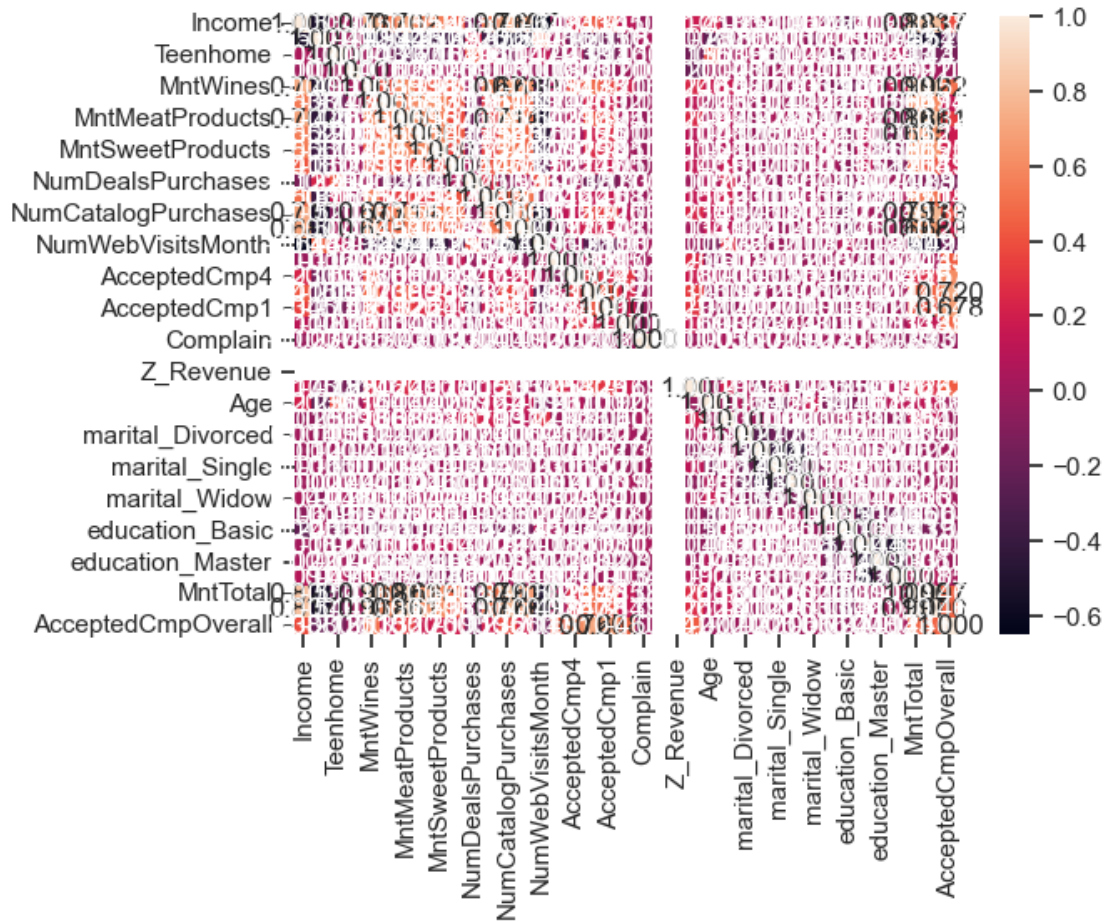
	AcceptedCmpOverall
0	0
1	0

```
2          0
3          0
4          0
```

```
[5 rows x 39 columns]
```

```
[ ]: sns.heatmap(data3.corr(), annot=True, fmt='.3f')
```

```
[ ]: <Axes: >
```



```
[ ]: # DataFrame
def make_corr_df(df):
    cr = data3.corr()
    cr = cr.abs().unstack()
    cr = cr.sort_values(ascending=False)
    cr = cr[cr >= 0.3]
    cr = cr[cr < 1]
    cr = pd.DataFrame(cr).reset_index()
```



```
cr.columns = ['f1', 'f2', 'corr']
return cr
```

```
[ ]: #
def corr_groups(cr):
    grouped_feature_list = []
    correlated_groups = []

    for feature in cr['f1'].unique():
        if feature not in grouped_feature_list:
            #
            correlated_block = cr[cr['f1'] == feature]
            cur_dups = list(correlated_block['f2'].unique()) + [feature]
            grouped_feature_list = grouped_feature_list + cur_dups
            correlated_groups.append(cur_dups)
    return correlated_groups
```

```
[ ]: #
corr_groups(make_corr_df(data3))
```

```
[ ]: [['MntTotal',
      'MntWines',
      'MntMeatProducts',
      'Income',
      'NumCatalogPurchases',
      'NumStorePurchases',
      'MntFishProducts',
      'MntSweetProducts',
      'MntFruits',
      'Kidhome',
      'NumWebPurchases',
      'NumWebVisitsMonth',
      'AcceptedCmp5',
      'AcceptedCmpOverall',
      'MntGoldProds',
      'AcceptedCmp1',
      'MntRegularProds'],
      ['AcceptedCmpOverall', 'MntWines', 'AcceptedCmp5', 'AcceptedCmp4'],
      ['education_Graduation', 'education_PhD'],
      ['marital_Married', 'marital_Single', 'marital_Together'],
      ['AcceptedCmpOverall', 'AcceptedCmp2'],
      ['education_Graduation', 'education_Master'],
      ['AcceptedCmpOverall', 'AcceptedCmp3'],
      ['AcceptedCmpOverall', 'AcceptedCmp5', 'Response'],
      ['Teenhome', 'NumWebVisitsMonth', 'NumDealsPurchases'],
      ['Teenhome', 'Age'],
      ['education_Graduation', 'education_2n Cycle']]
```

## 0.10

```
[ ]: X3_ALL = data3.drop(['Recency'], axis=1)

[ ]: #
X3_train, X3_test, y3_train, y3_test = train_test_split(X3_ALL,
↳data3['Recency'],
                                                    test_size=0.2,
                                                    random_state=1)

[ ]: #      L1-
e_lr1 = LogisticRegression(C=1000, solver='liblinear', penalty='l1',
↳max_iter=500, random_state=1)
e_lr1.fit(X3_train, y3_train)
#
e_lr1.coef_

[ ]: array([[ 2.36611132e-05, -4.37617142e-01, -1.06800247e-02, ...,
           -1.63170820e-03, -8.91019230e-05, -1.06604264e+00],
          [-4.08684338e-06,  2.11252986e-01,  3.26968614e-01, ...,
           2.06092631e-04, -5.67253475e-04, -3.79701869e-02],
          [ 2.97309082e-05,  2.78608379e-01,  7.84637108e-01, ...,
           7.60268573e-04,  9.22122307e-04, -2.72190745e-01],
          ...,
          [-3.33796349e-06,  5.87316491e-01,  8.31406329e-01, ...,
           1.86837335e-04,  9.46010783e-05,  2.65448813e-02],
          [-3.64459359e-05, -1.11298763e+00,  2.33999344e-01, ...,
          -2.42398207e-04,  1.69350949e-04, -6.53642523e-02],
          [ 4.67890738e-05, -1.67844004e+00,  7.82124718e-01, ...,
          -4.24294712e-05, -6.97529794e-04,  1.87736876e-02]])

[ ]: #      "      "
from sklearn.feature_selection import SelectFromModel
sel_e_lr1 = SelectFromModel(e_lr1)
sel_e_lr1.fit(X3_train, y3_train)
sel_e_lr1.get_support()

[ ]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
           True,  True,  True,  True,  True,  True,  True,  True,  True,
           True,  True,  True,  True,  True,  True,  True,  True,  True,
           True,  True])

[ ]: e_lr2 = LinearSVC(C=0.01, penalty="l1", max_iter=2000, dual=False)
e_lr2.fit(X3_train, y3_train)
#
e_lr2.coef_
```

/Users/seralekhin/BMSTU\_Labs/.env/lib/python3.11/site-

```
packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
```

```
warnings.warn(
```

```
[ ]: array([[ -8.75481522e-07,  0.00000000e+00,  0.00000000e+00, ...,
          -1.22148808e-05,  0.00000000e+00,  0.00000000e+00],
          [-1.70902922e-06,  0.00000000e+00,  0.00000000e+00, ...,
          -3.98938091e-05,  4.26048994e-05,  0.00000000e+00],
          [-3.54264966e-06,  0.00000000e+00,  0.00000000e+00, ...,
          -1.39468632e-04,  9.96158201e-05,  0.00000000e+00],
          ...,
          [-2.96785441e-06,  0.00000000e+00,  0.00000000e+00, ...,
           3.98224071e-05,  2.94380051e-04,  0.00000000e+00],
          [-3.02787186e-06,  0.00000000e+00,  0.00000000e+00, ...,
           1.29322152e-04,  0.00000000e+00,  0.00000000e+00],
          [ 4.46651443e-07,  0.00000000e+00,  0.00000000e+00, ...,
           0.00000000e+00, -1.38467919e-04,  0.00000000e+00]])
```

```
[ ]: # False . .
sel_e_lr2 = SelectFromModel(e_lr2)
sel_e_lr2.fit(X3_train, y3_train)
sel_e_lr2.get_support()
```

```
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-
```

```
packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
```

```
warnings.warn(
```

```
[ ]: array([ True, False, False,  True,  True,  True,  True,  True,  True,
           True,  True, False,  True, False, False, False, False, False,
          False, False, False,  True, False,  True,  True, False, False,
          False, False, False, False, False, False, False, False,  True,
           True, False])
```