

Московский государственный технический университет им. Н.Э. Баумана
Факультет «Информатика и системы управления»
Кафедра «Системы обработки информации и управления»



Лабораторная работа №4
по дисциплине «Методы машинного обучения»
«Алгоритм Policy Iteration»

ИСПОЛНИТЕЛЬ:

Алехин С.С.
Группа ИУ5-23М

ПРОВЕРИЛ:

Балашов А.М.

Задание

1. На основе рассмотренного на лекции примера реализуйте алгоритм Policy Iteration для любой среды обучения с подкреплением (кроме рассмотренной на лекции среды Toy Text / Frozen Lake) из библиотеки [Gym](#) (или аналогичной библиотеки).

Lab4

June 21, 2023

```
[ ]: ! pip install gymnasium
      ! pip install pygame

import os
import pygame
import gymnasium as gym
import numpy as np
from pprint import pprint
```

```
Requirement already satisfied: gymnasium in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (0.28.1)
Requirement already satisfied: numpy>=1.21.0 in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from gymnasium)
(1.25.0)
Requirement already satisfied: jax-jumpy>=1.0.0 in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from gymnasium)
(1.0.0)
Requirement already satisfied: cloudpickle>=1.2.0 in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from gymnasium)
(2.2.1)
Requirement already satisfied: typing-extensions>=4.3.0 in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from gymnasium)
(4.6.3)
Requirement already satisfied: farama-notifications>=0.0.1 in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (from gymnasium)
(0.0.4)
Requirement already satisfied: pygame in
/Users/seralekhin/BMSTU_Labs/.env/lib/python3.11/site-packages (2.4.0)
```

```
[ ]: class PolicyIterationAgent:
      """
      """
      def __init__(self, env):
          self.env = env
          #
          self.observation_dim = 4 * 12
          #
```

```

# https://gymnasium.farama.org/environments/toy_text/taxi/
self.actions_variants = np.array([0,1,2,3])
# 0: Move south (down)
# 1: Move north (up)
# 2: Move east (right)
# 3: Move west (left)
# 4: Pickup passenger
# 5: Drop off passenger
#         (      )
#    4 12  4
self.policy_probs = np.full((self.observation_dim, len(self.
↪actions_variants)), 0.25)
#           v(s)
self.state_values = np.zeros(shape=(self.observation_dim))
#
self.maxNumberOfIterations = 1000
self.theta=1e-6
self.gamma=0.99

def print_policy(self):
    '''

    '''
    print('      :')
    pprint(self.policy_probs)

def policy_evaluation(self):
    '''

    '''
    #
    valueFunctionVector = self.state_values
    for iterations in range(self.maxNumberOfIterations):
        #
        valueFunctionVectorNextIteration=np.zeros(shape=(self.
↪observation_dim))
        #
        for state in range(self.observation_dim):
            #
            action_probabilities = self.policy_probs[state]
            #
            outerSum=0
            for action, prob in enumerate(action_probabilities):
                innerSum=0
                #
                for probability, next_state, reward, isTerminalState in_
↪self.env.P[state][action]:

```

```

        innerSum=innerSum+probability*(reward+self.gamma*self.
↪state_values[next_state])
        outerSum=outerSum+self.policy_probs[state][action]*innerSum
        valueFunctionVectorNextIteration[state]=outerSum
        if(np.max(np.
↪abs(valueFunctionVectorNextIteration-valueFunctionVector))<self.theta):
            #
            valueFunctionVector=valueFunctionVectorNextIteration
            break
        valueFunctionVector=valueFunctionVectorNextIteration
    return valueFunctionVector

def policy_improvement(self):
    """
    """

    qvaluesMatrix=np.zeros((self.observation_dim, len(self.
↪actions_variants)))
    improvedPolicy=np.zeros((self.observation_dim, len(self.
↪actions_variants)))
    #
    for state in range(self.observation_dim):
        for action in range(len(self.actions_variants)):
            for probability, next_state, reward, isTerminalState in self.
↪env.P[state][action]:
                ↵
↪qvaluesMatrix[state,action]=qvaluesMatrix[state,action]+probability*(reward+self.
↪gamma*self.state_values[next_state])

            #
            bestActionIndex=np.where(qvaluesMatrix[state,:]==np.
↪max(qvaluesMatrix[state,:]))
            #
            improvedPolicy[state,bestActionIndex]=1/np.size(bestActionIndex)
    return improvedPolicy

def policy_iteration(self, cnt):
    """
    """

    policy_stable = False
    for i in range(1, cnt+1):
        self.state_values = self.policy_evaluation()
        self.policy_probs = self.policy_improvement()

```

```
print(f'          {i}      .')
```

```
[ ]: def play_agent(agent):
    env2 = gym.make('CliffWalking-v0', render_mode='human')
    state = env2.reset()[0]
    done = False
    while not done:
        p = agent.policy_probs[state]
        if isinstance(p, np.ndarray):
            action = np.random.choice(len(agent.actions_variants), p=p)
        else:
            action = p
        next_state, reward, terminated, truncated, _ = env2.step(action)
        env2.render()
        state = next_state
        if terminated or truncated:
            done = True
```

```
[ ]: #
env = gym.make('CliffWalking-v0')
env.reset()
#
agent = PolicyIterationAgent(env)
agent.policy_iteration(1000)
agent.print_policy()
#
play_agent(agent)
```

1000 .

```
:
array([[0.          , 0.5          , 0.5          , 0.          ],
       [0.33333333, 0.33333333, 0.33333333, 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.33333333, 0.          , 0.33333333, 0.33333333],
       [0.          , 0.          , 0.5          , 0.5          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.5          , 0.5          , 0.          ],
       [0.          , 0.5          , 0.5          , 0.          ],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
```

```

[0.          , 0.33333333, 0.33333333, 0.33333333],
[0.          , 0.33333333, 0.33333333, 0.33333333],
[0.          , 0.33333333, 0.33333333, 0.33333333],
[0.          , 0.          , 0.5          , 0.5          ],
[0.          , 0.          , 0.5          , 0.5          ],
[0.          , 0.          , 1.          , 0.          ],
[0.          , 0.33333333, 0.33333333, 0.33333333],
[0.          , 0.5          , 0.          , 0.5          ],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.33333333, 0.33333333, 0.          , 0.33333333],
[0.          , 0.5          , 0.          , 0.5          ],
[0.          , 0.33333333, 0.33333333, 0.33333333],
[0.33333333, 0.          , 0.33333333, 0.33333333],
[0.5          , 0.          , 0.          , 0.5          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[1.          , 0.          , 0.          , 0.          ],
[0.5          , 0.5          , 0.          , 0.          ],
[0.33333333, 0.33333333, 0.33333333, 0.          ]]

```

Lab4

June 22, 2023

```
[ ]: import gymnasium as gym
import numpy as np
from pprint import pprint
```

```
[ ]: class PolicyIterationAgent:
    """
    """
    def __init__(self, env):
        self.env = env
        #
        self.observation_dim = 4 * 12
        #
        # https://gymnasium.farama.org/environments/toy\_text/taxi/
        self.actions_variants = np.array([0,1,2,3])
        # 0: Move south (down)
        # 1: Move north (up)
        # 2: Move east (right)
        # 3: Move west (left)
        # 4: Pickup passenger
        # 5: Drop off passenger
        # ( )
        # 4 12 4
        self.policy_probs = np.full((self.observation_dim, len(self.
→actions_variants)), 0.25)
        #  $v(s)$ 
        self.state_values = np.zeros(shape=(self.observation_dim))
        #
        self.maxNumberOfIterations = 1000
        self.theta=1e-6
        self.gamma=0.99

    def print_policy(self):
        """
        """
        print(' :')
```



```

pprint(self.policy_probs)

def policy_evaluation(self):
    """
    """
    #
    valueFunctionVector = self.state_values
    for iterations in range(self.maxNumberOfIterations):
        #
        valueFunctionVectorNextIteration=np.zeros(shape=(self.
↪observation_dim))
        #
        for state in range(self.observation_dim):
            #
            action_probabilities = self.policy_probs[state]
            #
            outerSum=0
            for action, prob in enumerate(action_probabilities):
                innerSum=0
                #
                for probability, next_state, reward, isTerminalState in_
↪self.env.P[state][action]:
                    innerSum=innerSum+probability*(reward+self.gamma*self.
↪state_values[next_state])
                    outerSum=outerSum+self.policy_probs[state][action]*innerSum
                valueFunctionVectorNextIteration[state]=outerSum
            if(np.max(np.
↪abs(valueFunctionVectorNextIteration-valueFunctionVector))<self.theta):
                #
                valueFunctionVector=valueFunctionVectorNextIteration
                break
        valueFunctionVector=valueFunctionVectorNextIteration
    return valueFunctionVector

def policy_improvement(self):
    """
    """
    qvaluesMatrix=np.zeros((self.observation_dim, len(self.
↪actions_variants)))
    improvedPolicy=np.zeros((self.observation_dim, len(self.
↪actions_variants)))
    #
    for state in range(self.observation_dim):

```

```

        for action in range(len(self.actions_variants)):
            for probability, next_state, reward, isTerminalState in self.
↳env.P[state][action]:

                ↳
↳qvaluesMatrix[state,action]=qvaluesMatrix[state,action]+probability*(reward+self.
↳gamma*self.state_values[next_state])

            #
            bestActionIndex=np.where(qvaluesMatrix[state,:]==np.
↳max(qvaluesMatrix[state,:]))
            #
            improvedPolicy[state,bestActionIndex]=1/np.size(bestActionIndex)
        return improvedPolicy

def policy_iteration(self, cnt):
    '''

    '''
    policy_stable = False
    for i in range(1, cnt+1):
        self.state_values = self.policy_evaluation()
        self.policy_probs = self.policy_improvement()
    print(f'          {i}      .')

```

```

[ ]: def play_agent(agent):
    env2 = gym.make('CliffWalking-v0', render_mode='human')
    state = env2.reset()[0]
    done = False
    while not done:
        p = agent.policy_probs[state]
        if isinstance(p, np.ndarray):
            action = np.random.choice(len(agent.actions_variants), p=p)
        else:
            action = p
        next_state, reward, terminated, truncated, _ = env2.step(action)
        env2.render()
        state = next_state
        if terminated or truncated:
            done = True

```

```

[ ]: #
env = gym.make('CliffWalking-v0')
env.reset()
#
agent = PolicyIterationAgent(env)
agent.policy_iteration(1000)

```

```
agent.print_policy()
#
play_agent(agent)
```

```
1000 .

:
array([[0.          , 0.5          , 0.5          , 0.          ],
       [0.33333333, 0.33333333, 0.33333333, 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.33333333, 0.          , 0.33333333, 0.33333333],
       [0.          , 0.          , 0.5          , 0.5          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.5          , 0.5          , 0.          ],
       [0.          , 0.5          , 0.5          , 0.          ],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.          , 0.5          , 0.5          ],
       [0.          , 0.          , 0.5          , 0.5          ],
       [0.          , 0.          , 1.          , 0.          ],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.          , 0.5          , 0.          , 0.5          ],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.33333333, 0.33333333, 0.          , 0.33333333],
       [0.          , 0.5          , 0.          , 0.5          ],
       [0.          , 0.33333333, 0.33333333, 0.33333333],
       [0.33333333, 0.          , 0.33333333, 0.33333333],
       [0.5          , 0.          , 0.          , 0.5          ],
       [1.          , 0.          , 0.          , 0.          ],
       [1.          , 0.          , 0.          , 0.          ],
       [1.          , 0.          , 0.          , 0.          ],
       [1.          , 0.          , 0.          , 0.          ]],
```

```
[1.      , 0.      , 0.      , 0.      ],  
[1.      , 0.      , 0.      , 0.      ],  
[1.      , 0.      , 0.      , 0.      ],  
[1.      , 0.      , 0.      , 0.      ],  
[0.5     , 0.5     , 0.      , 0.      ],  
[0.33333333, 0.33333333, 0.33333333, 0.      ]])
```

The Kernel crashed while executing code in the the current cell or a previous cell. Please review the code in the cell(s) to identify a possible cause of the failure. Click [here](https://aka.ms/vscodeJupyterKernelCrash) for more info. View Jupyter [log](command:jupyter.viewOutput) for further details.